


Lab 7

K-means Clustering

Let's make up some data to cluster.

```
tmp <- c(rnorm(30, -3), rnorm(30, 3))  
x <- cbind(x= tmp, y= rev(tmp))  
plot(x)
```



The function to do k-means clustering in base R is called 'kmeans()'.

```
kmeans(x, centers=2)
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

```
      x      y
1  3.083912 -2.919432
2 -2.919432  3.083912
```

Clustering vector:

```
[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```

Within cluster sum of squares by cluster:

```
[1] 49.82914 49.82914
(between_SS / total_SS =  91.6 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"       "
```

```
km <- kmeans(x, centers = 4, nstart=20)
km
```

K-means clustering with 4 clusters of sizes 18, 18, 12, 12

Cluster means:

```
      x      y
1 -3.493336  2.716752
2  2.716752 -3.493336
3  3.634653 -2.058575
4 -2.058575  3.634653
```

Clustering vector:

```
[1] 4 1 4 1 1 4 1 4 1 1 4 4 4 1 1 4 1 1 4 1 1 1 4 1 1 1 1 4 1 1 4 3 2 3 2 2 2 2 3
[39] 2 2 2 3 2 2 3 2 2 3 3 3 2 2 3 2 3 2 2 3 2 3 2 3 2 3 2 3 2 3 2 3 2 3 2 3 2 3
```

Within cluster sum of squares by cluster:

```
[1] 15.14244 15.14244 13.79891 13.79891
```

```
(between_SS / total_SS = 95.1 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

##Q. What 'component' of your results object details - cluster size - cluster assignment/membership - cluster center

```
#will provide cluster size
km$size
```

```
[1] 18 18 12 12
```

```
#will provide cluster center
km$center
```


```
      x      y
1 -3.493336  2.716752
2  2.716752 -3.493336
3  3.634653 -2.058575
4 -2.058575  3.634653
```

```
#Membership vector
km$cluster
```

```
[1] 4 1 4 1 1 4 1 4 1 1 4 4 4 1 1 4 1 1 1 4 1 1 1 1 4 1 1 1 1 4 1 4 3 2 3 2 2 2 2 3
[39] 2 2 2 3 2 2 3 2 2 3 3 3 2 2 3 2 3 2 2 3 2 3
```

##Q. Plot x colored by the kmeans cluster assignment and add cluster centers as blue points

```
plot(x, col=km$cluster)
points(km$centers, col="blue", pch=15, cex=1.5)
```



hclust()

Hierarchical Clustering

```
hc <- hclust( dist(x))
hc
```

Call:

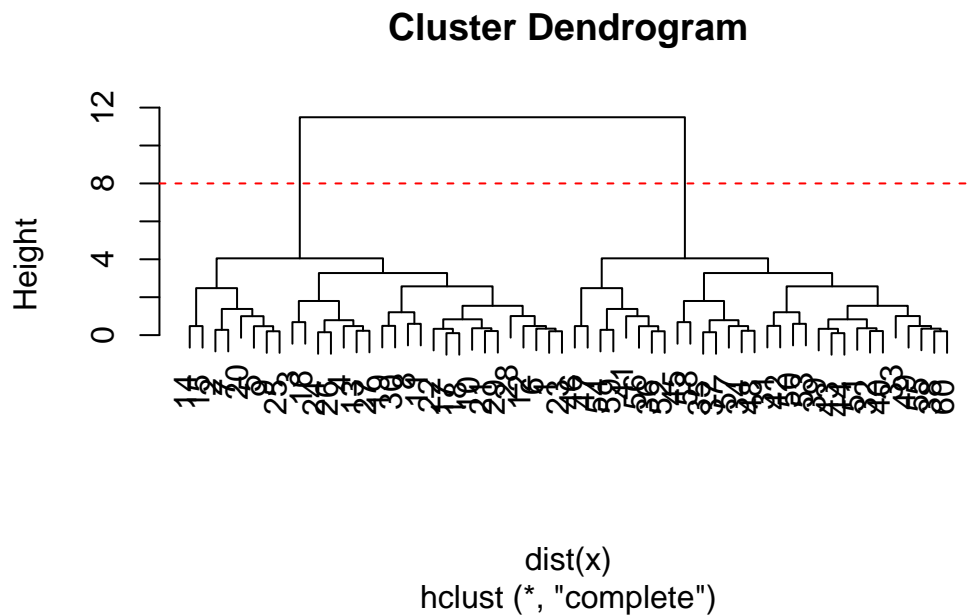
```
hclust(d = dist(x))
```

Cluster method : complete

Distance : euclidean

Number of objects: 60

```
plot(hc)
abline(h=8, col="red", lty=2)
```



To get my “main” result (cluster membership), I want to

```
cutree(hc, h=8)
```


```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2
[39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

More often we will use ‘cutree()’ with k=2 for example

```
grps = cutree(hc, k=3)
```

Make a plot of our ‘hclust()’ results i.e. our data colored by cluster assignment

```
plot(x, col= grps)
```



Principal Component Analysis (PCA)

```
url <- "https://tinyurl.com/UK-foods"  
x <- read.csv(url)
```

##Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

```
nrow(x)
```

```
[1] 17
```

```
ncol(x)
```

```
[1] 5
```

There are 17 rows and 5 columns.

```
View(x)
head(x)
```

	X	England	Wales	Scotland	N.Ireland
1	Cheese	105	103	103	66
2	Carcass_meat	245	227	242	267
3	Other_meat	685	803	750	586
4	Fish	147	160	122	93
5	Fats_and_oils	193	235	184	209
6	Sugars	156	175	147	139

```
tail(x)
```

	X	England	Wales	Scotland	N.Ireland
12	Fresh_fruit	1102	1137	957	674
13	Cereals	1472	1582	1462	1494
14	Beverages	57	73	53	47
15	Soft_drinks	1374	1256	1572	1506
16	Alcoholic_drinks	375	475	458	135
17	Confectionery	54	64	62	41

```
rownames(x) <- x[,1]
x <- x[,-1]
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

```
dim(x)
```

```
[1] 17 4
```


```
x <- read.csv(url, row.names=1)
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

##Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

The second approach (`x <- read.csv(url, row.names=1)`) because it is a single step and easier to remember. They are both equally robust if the dataset is the same.


```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```



##Q3: Changing what optional argument in the above `barplot()` function results in the following plot?

Changing `beside=TRUE` to `beside=FALSE` will provide the following plot.

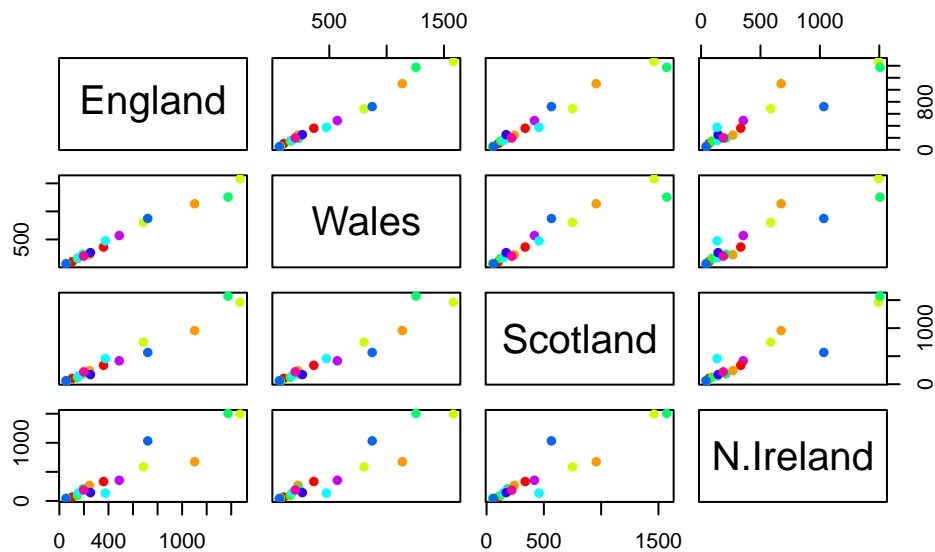
```
#changing beside=TRUE to beside=FALSE will provide the following plot
barplot(as.matrix(x), beside=FALSE, col=rainbow(nrow(x)))
```



##Q5: Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

Yes, we are generating pairwise plots by using `pairs()` using X dataset and indicating that the colors of each datapoint will be based off the rainbow colors (multiple colors instead of one color) and the `pch=16` just sets the shape of the datapoint on the plot. The resulting pairwise plots provides plots of each column of the dataset (x) against each other (countries) and the named countries in the empty boxes indicate the y-axis and x-axis of each plot. For example all the plots in the England column has England's data as the x-axis and the plots in the England row has England's data as the y-axis. If a given point lies on the diagonal for a given plot it means that the value is the same for the two countries (x-axis and y-axis).

```
pairs(x, col=rainbow(10), pch=16)
```



PCA to the rescue!

The main function in base R to do PCA is called 'prcomp()'. One issue with the 'prcomp()' function is that it expects the transpose of our data as input.

##Q6. What is the main differences between N. Ireland and the other countries of the UK in terms of this data-set? A main difference between N. Ireland and other countries is that the data has some outliers more outliers compared to the rest.


```
# Use the prcomp() PCA function to return a list object.
pca <- prcomp( t(x) )
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	4.189e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00


##Q7. Complete the code below to generate a plot of PC1 vs PC2. The second line adds text labels over the data points.

```
# Plot PC1 vs PC2
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2", xlim=c(-270,500))
text(pca$x[,1], pca$x[,2], colnames(x))
```



##Q8. Customize your plot so that the colors of the country names match the colors in our UK and Ireland map and table at start of this document.

```
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2", xlim=c(-270,500))
text(pca$x[,1], pca$x[,2], colnames(x), col=c("orange", "red", "blue", "darkgreen"))
```



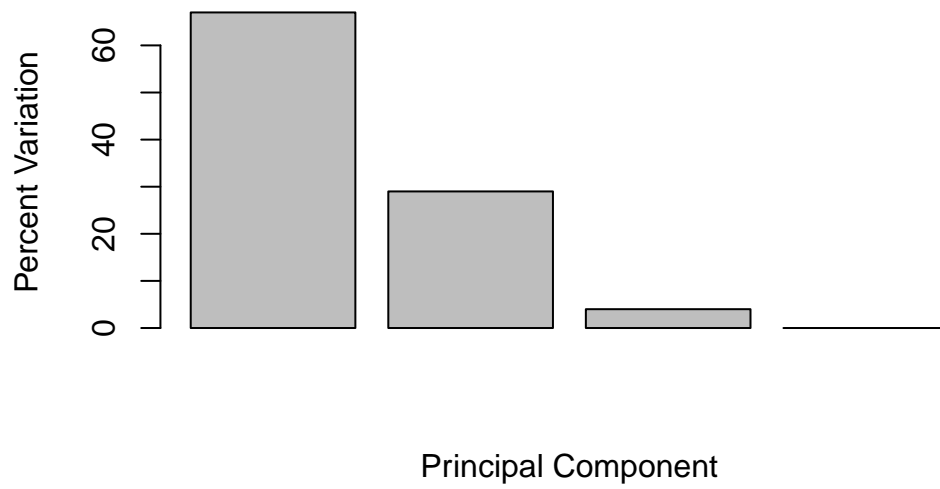
```
#Code below: Use square of pca$sdev , which stands for "standard deviation", to calculate
v <- round( pca$sdev^2/sum(pca$sdev^2) * 100 )
v
```

```
[1] 67 29 4 0
```


```
#Alternative method to calculate how much variation in the original data
z <- summary(pca)
z$importance
```

	PC1	PC2	PC3	PC4
Standard deviation	324.15019	212.74780	73.87622	4.188568e-14
Proportion of Variance	0.67444	0.29052	0.03503	0.000000e+00
Cumulative Proportion	0.67444	0.96497	1.00000	1.000000e+00

```
barplot(v, xlab="Principal Component", ylab="Percent Variation")
```



```
## Lets focus on PC1 as it accounts for > 90% of variance  
par(mar=c(10, 3, 0.35, 0))  
barplot( pca$rotation[,1], las=2 )
```



##Q9: Generate a similar 'loadings plot' for PC2. What two food groups feature prominently and what does PC2 mainly tell us about?

Fresh potatoes and Soft Drinks. PC2 plot is explaining how each variable is contributing/affecting PC2.

```
par(mar=c(10, 3, 0.35, 0))
barplot( pca$rotation[,2], las=2 )
```

